

Transport Reliability Analytics

Technical Report

Land Transport Authority (LTA)

Singapore MRT & Bus Network

Stack: Go, gRPC, PostgreSQL, Next.js
Deployment: Docker Compose
Performance: 100+ QPS with 95%+ success rate

October 19, 2025

Contents

1	Executive Summary	2
1.1	Quick Start	2
2	System Architecture	2
2.1	Component Diagram	2
2.2	Technology Stack	2
3	Getting Started	2
3.1	Prerequisites	2
3.2	Installation	2
3.3	Access Points	3
4	API Documentation	3
4.1	Endpoints	3
4.2	Example: Create Incident	3
4.3	Example: Get Top Breakdowns	3
5	Database Schema	3
5.1	Tables	3
5.2	Performance Indexes	3
5.3	Seed Data	4
6	Performance Testing	4
6.1	Stress Test Results	4
6.2	Running Tests	4
7	Frontend Dashboard	4
7.1	Features	4
7.2	Technology	5
8	Production Considerations	5
8.1	Scaling	5
8.2	Monitoring	5
8.3	Security	5
9	Troubleshooting	5
9.1	Common Issues	5
10	Conclusion	6
10.1	Project Structure	6

1 Executive Summary

This document provides an overview of the Transport Reliability Analytics system - a production-ready full-stack application for tracking and analyzing transport incidents across Singapore's MRT and bus network.

1.1 Quick Start

Start the entire system with one command:

```
1 docker-compose up
```

Access the dashboard at: <http://localhost:3000>

2 System Architecture

2.1 Component Diagram




Figure 1: System Architecture

2.2 Technology Stack

Backend: Go 1.24, go-coldbrew, gRPC, PostgreSQL 15

Frontend: Next.js 15, React 18, TypeScript, Tailwind CSS

Infrastructure: Docker, nginx

3 Getting Started

3.1 Prerequisites

- Docker and Docker Compose
- Ports 3000, 8080, 9090, 9091, 5433 available

3.2 Installation

Automated Setup:

```
1 ./setup.sh
```

Manual Setup:

```
1 docker-compose up
```

3.3 Access Points

Service	URL	Port
Dashboard	http://localhost:3000	3000
API	http://localhost:8080	8080
Backend	http://localhost:9091	9091
Database	localhost:5433	5433

4 API Documentation

4.1 Endpoints

Method	Endpoint	Description
POST	/incidents	Submit new incident
GET	/analytics/top_breakdowns	Top N lines/stations
GET	/analytics/mean_time_between_failures	MTBF calculation
GET	/analytics/recent_disruptions	Recent incidents

4.2 Example: Create Incident

```

1 curl -X POST http://localhost:8080/incidents \
2   -H "Content-Type: application/json" \
3   -d '{
4     "line": "North South Line",
5     "station": "Orchard",
6     "timestamp": "2025-10-16T08:32:00Z",
7     "duration_minutes": 45,
8     "incident_type": "signal"
9   }'
```

4.3 Example: Get Top Breakdowns

```

1 curl "http://localhost:8080/analytics/top_breakdowns?scope=station&limit=5"
```

5 Database Schema

5.1 Tables

lines: id (UUID), name (TEXT), created_at

stations: id (UUID), name (TEXT), line_id (FK), status, created_at

incidents: id (UUID), station_id (FK), line_id (FK), ts, duration_minutes, incident_type, status

5.2 Performance Indexes

- `idx_incidents_ts` - Timestamp descending
- `idx_incidents_line_ts` - Line + timestamp (covering index)
- `idx_incidents_station_ts` - Station + timestamp (covering index)
- `idx_incidents_status` - Status filtering

5.3 Seed Data

- 6 MRT lines (NSL, EWL, CCL, DTL, TEL, NEL)
- 150+ stations
- 400+ incidents (last 90 days)

6 Performance Testing

6.1 Stress Test Results

Two load tests were performed to validate system performance under real conditions:

Light Load Test (10 QPS):

Metric	Result
Total Requests	600
Success Rate	100%
Mean Latency	5.2ms
p50 Latency	5.1ms
p95 Latency	9.4ms
p99 Latency	11.2ms

Heavy Load Test (100 QPS):

Metric	Result
Total Requests	6,000
Success Rate	100%
Mean Latency	3.0ms
p50 Latency	2.6ms
p95 Latency	6.0ms
p99 Latency	8.0ms

6.2 Running Tests

```
1 # Run all stress tests
2 make stress-test
3
4 # Individual tests
5 make stress-test-10qps
6 make stress-test-100qps
```

7 Frontend Dashboard

7.1 Features

- Auto-refresh every 30 seconds
- Interactive bar charts (top breakdowns)
- Color-coded MTBF metrics

- Sortable disruptions table
- Manual refresh button

7.2 Technology

Next.js 15, TypeScript, Tailwind CSS, Recharts

8 Production Considerations

8.1 Scaling

- Horizontal: Multiple app instances behind load balancer
- Database: Read replicas for analytics
- Connection pooling: 25 max, 5 idle

8.2 Monitoring

- Health checks: `/health` and `/ready`
- Structured logging with request IDs
- Prometheus metrics endpoint

8.3 Security

- Change default passwords
- Enable SSL/TLS
- Implement API authentication
- Add rate limiting
- CORS via nginx

9 Troubleshooting

9.1 Common Issues

Port Conflicts: Edit `docker-compose.yml` to change ports

Database Issues:

```
1 docker-compose logs db
2 docker-compose exec db psql -U lta_user -d transport_reliability
```

Frontend API Errors:

1. Verify nginx: `docker-compose ps nginx`
2. Test API: `curl http://localhost:8080/health`
3. Hard refresh browser: `Cmd+Shift+R`

10 Conclusion

The Transport Reliability Analytics system is a production-ready application demonstrating:

- Modern microservices architecture
- High-performance API (100+ QPS)
- Real-time data visualization
- Professional Docker deployment

10.1 Project Structure

```
1 .
2     backend/          # Go services
3     database/         # PostgreSQL setup
4     frontend/         # Next.js dashboard
5     nginx/            # Reverse proxy
6     proto/            # Protocol buffers
7     scripts/          # Stress tests
8     docker-compose.yml
```

Repository: [stunning-octo-eureka](#)

Tech Stack: Go · gRPC · PostgreSQL · Next.js · Docker · nginx